ParticleGAN 0.9.5: calibration convergence

Native bispf_optimized; continuous domain | fold 0, seed 0 | G 0.006, D 0.009, particles 0.015

Latest saved update 20,000 / 20,000 | audited complete calibration curve

Donor energy score (lower is better)

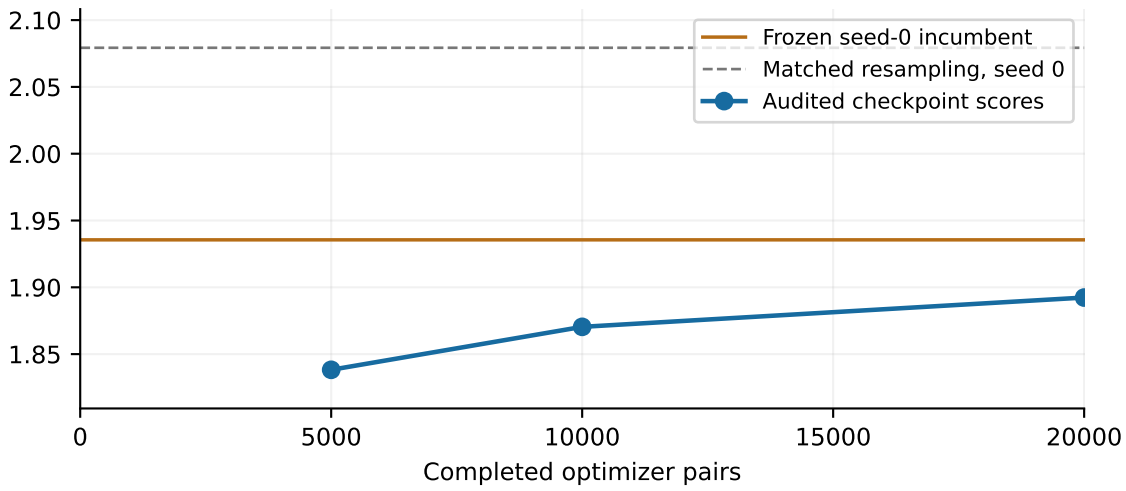

Pooled cell energy distance (lower is better)

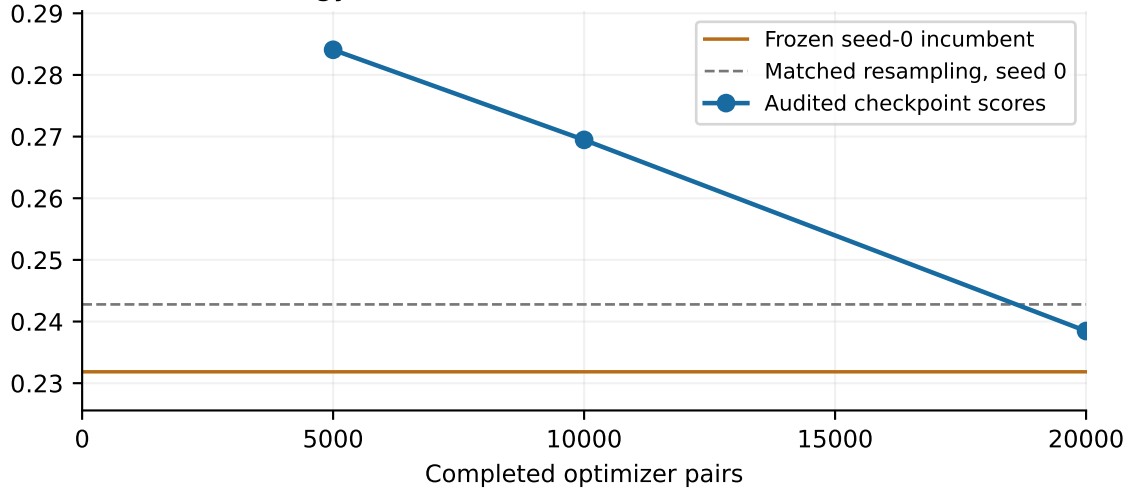

Native adversarial losses

Observed gradient norms

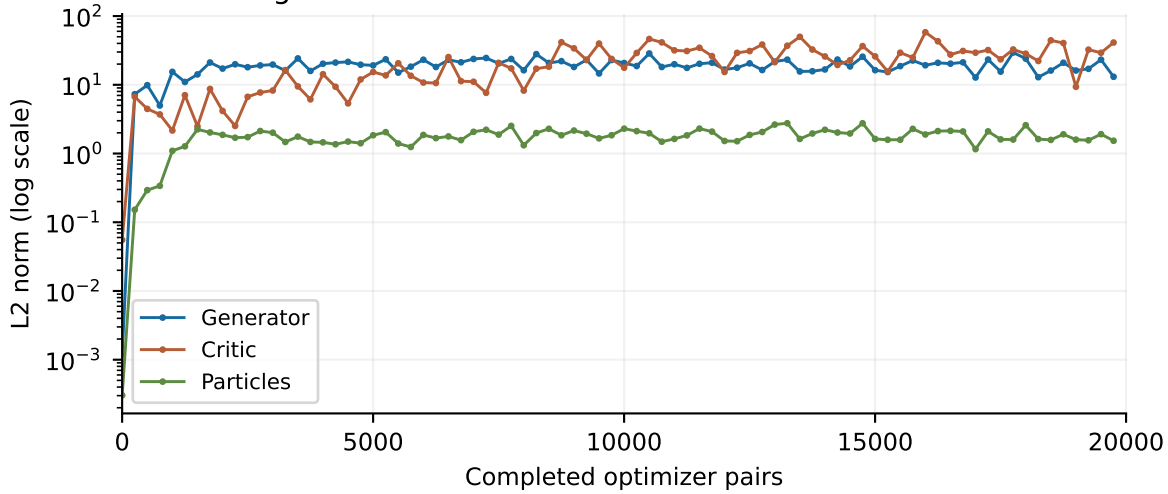

Latent particle covariance trace

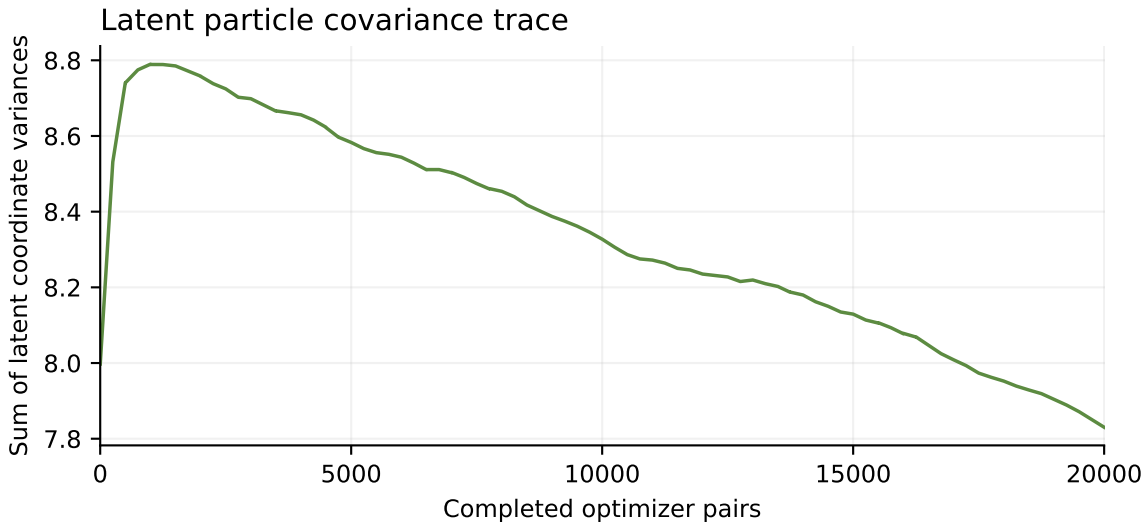

Actual native optimizer learning rates

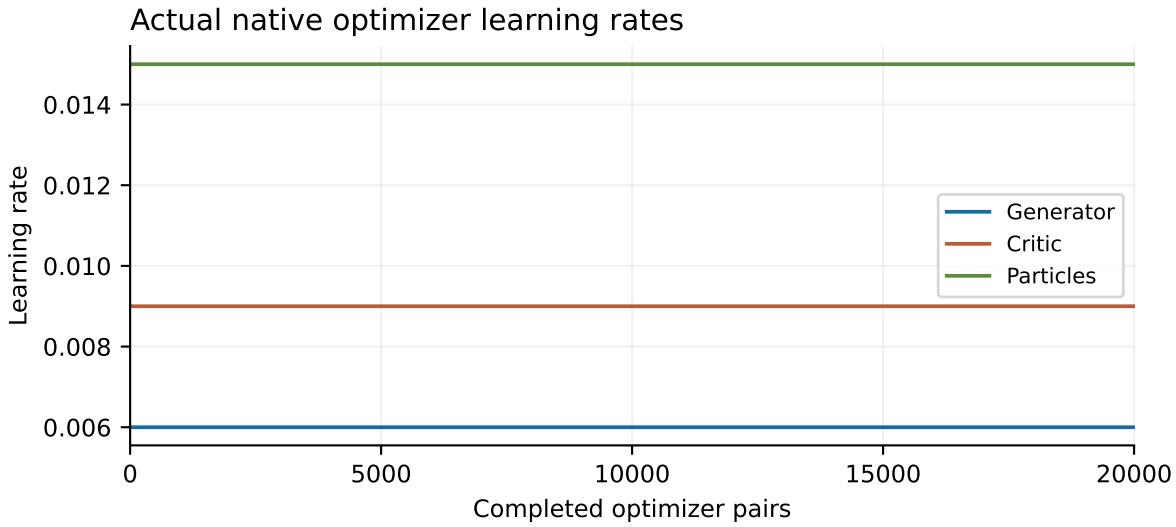

Fixed 20,000-update endpoint selects the model. The 5k and 10k checkpoints are convergence diagnostics.

Four adaptive calibration donors; two generation replicates per checkpoint. Evaluation donors are separate.