Released ParticleGAN: does longer training improve the best rate setting?

Same model and seed | generator LR 0.006, critic LR 0.009, particle LR 0.015

Parent 20k results informed the duration test. Only the declared 40k extension endpoint selects it.

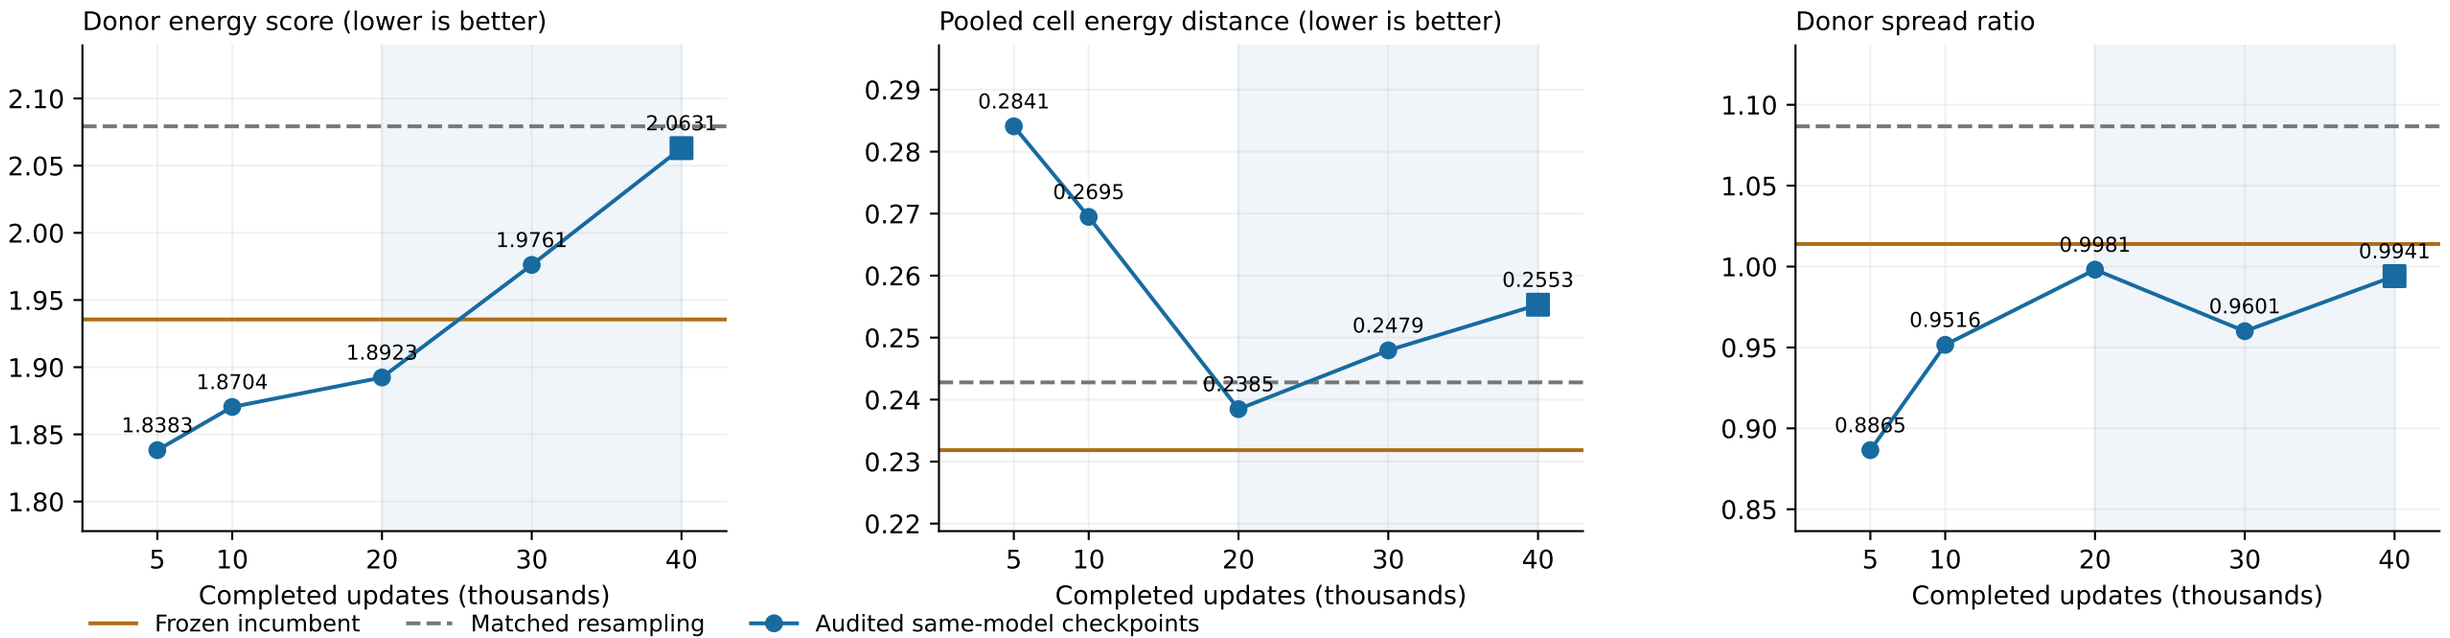

Four adaptive calibration donors; two generation replicates per checkpoint. Evaluation donors excluded.

Shading: new 20k→40k phase. Square: fixed 40k endpoint. Other points retain their original diagnostic/historical roles.