Generator rate: matched convergence comparison

Released bispf_optimized | fixed D 0.009, particles 0.015 | fold 0, seed 0

Same four calibration donors; two generation replicas per checkpoint. Both runs complete 20,000 updates.

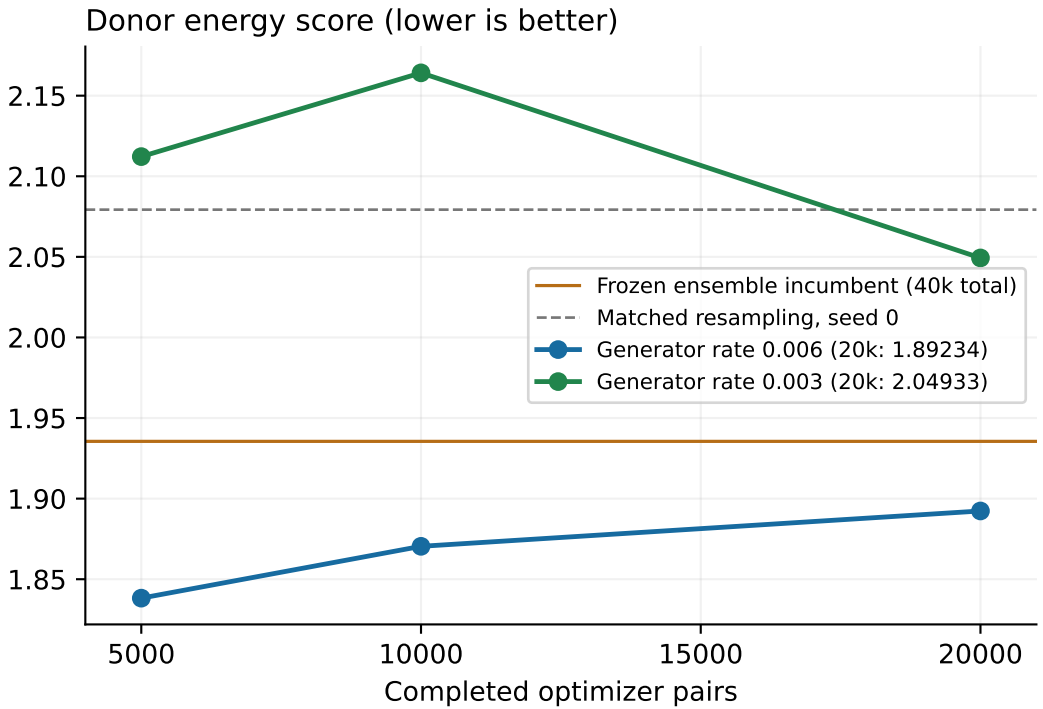

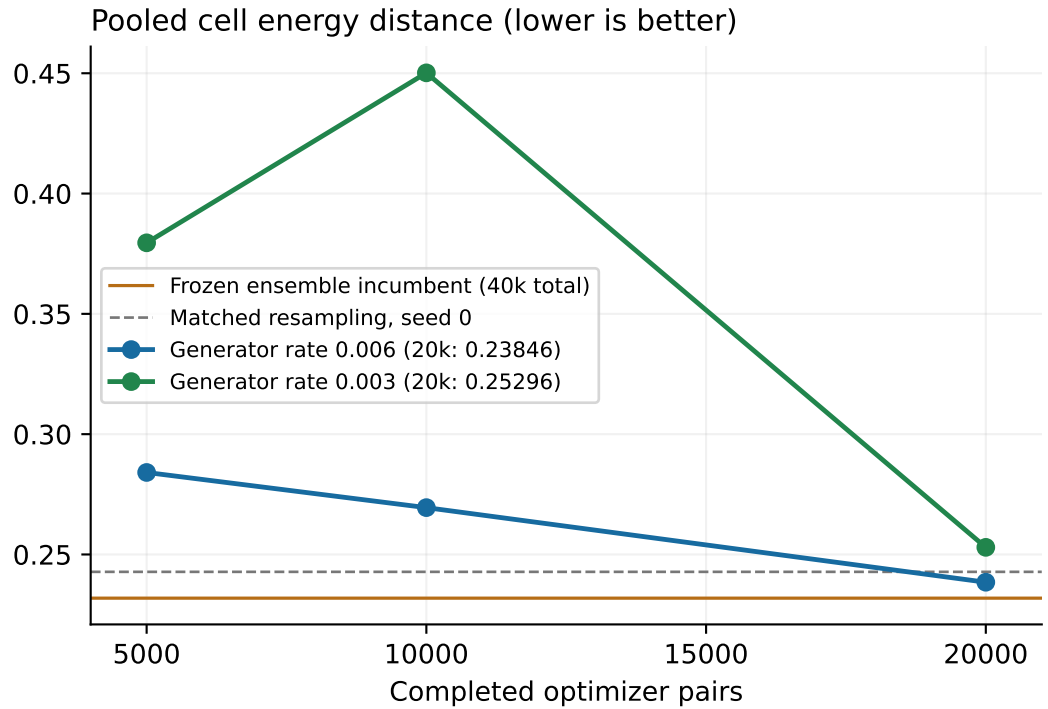

Only the final 20k endpoint can select a model; 5k and 10k scores diagnose convergence.

Adaptive calibration, one training seed. These curves do not establish a stable joint ATH or evaluation performance.

Parameter-dependent random draws may diverge despite fixed seed streams. The 20 evaluation donors remain excluded.